

AION Seven-Day Open-Mission Compounding Programme

Frozen Protocol, Matched Ablations, Delayed Authority, Repair, Retention, Operations,
and Verdict Criteria

AION / COMDEX Technical Record

Campaign initiated 8 August 2026
Document snapshot: 8 August 2026, Europe/Madrid

Contents

1 Purpose and Research Question	3
2 Frozen Campaign Contract	3
3 What the Five Experimental Arms Mean	4
4 Mission Contract and Lifecycle	4
5 Implemented Mission Authorities	5
5.1 Answer-hidden polyglot systems fault	5
5.2 Rolling later-public-change authority	6
5.3 Precommitted repository-repair mission	6
6 Retention Protocol	6
7 Repair Protocol	7
8 Metrics and Claim Gates	7
8.1 Live baseline at document creation	7
9 Authority Separation and Safety	8
10 Operational Services and Reboot Persistence	8
11 Focused Dashboard	9
12 Expected Seven-Day Progression	9
13 End-State Verdicts	10
13.1 Minimum operational pass	10
13.2 “OK” intelligence result	10
13.3 “Fucking wow” result	11
13.4 Flat or negative result	11
14 Known Limitations and Required Corrections	11
15 Artifact and Evidence Map	12
16 Verification and Continuation Procedure	13

1 Purpose and Research Question

The seven-day open-mission programme marks a deliberate change in AION’s development strategy. Technical AGA v1 is frozen as a completed technical baseline. The programme no longer treats additional internal project counts, curriculum receipts, or repeated passes in the same task family as evidence of increasing generality. Its purpose is to test whether AION’s governed memory, repair, retention, authority separation, and accumulated experience produce a measurable advantage over the same underlying proposer without those mechanisms.

The primary research question is:

After controlling the proposer, tools, and action budget, does full AION solve unfamiliar consequential missions more successfully or more efficiently than proposer-only, cold, no-memory, and no-repair controls; and are any gains retained after sources close and confirmed by later independent outcomes?

The intended compounding chain is

$$\begin{aligned} \text{unfamiliar mission} &\rightarrow \text{frozen success contract} \rightarrow \text{matched five-arm execution} \\ &\rightarrow \text{independent later consequence} \rightarrow \text{failure attribution} \\ &\rightarrow \text{governed repair or abstention} \rightarrow \text{source closure} \\ &\rightarrow \text{seven-day protected transfer} \rightarrow \text{harder next mission.} \end{aligned} \tag{1}$$

This is a path-to-evidence experiment, not an automatic AGI declaration. The campaign contract explicitly fixes `agi_claim_authorized` to false.

2 Frozen Campaign Contract

The governor created an immutable campaign contract before matched execution. The live contract is summarized in Table 1.

Table 1: Frozen seven-day campaign contract.

Field	Frozen value
Campaign identifier	<code>aion_seven_day_open_mission_moonshot_202608</code>
Schema	<code>aion.hexcore.open_mission_compounding.v1</code>
Objective	Measure whether governed accumulated experience improves verified performance on unfamiliar real missions under declining scaffolding.
Start	8 August 2026, 19:26:28 UTC (21:26:28 CEST)
Nominal end	15 August 2026, 19:26:28 UTC (21:26:28 CEST)
Duration	604,800 seconds (seven days)
Fixed proposer	<code>ollama:qwen3:1.7b@local</code>
Action budget	40 actions per arm per cohort
Contract SHA-256	<code>ddddbea7ac14ce13f288629d31050e6103e82a373053c17755f42ba0c789cd4</code>
Tool-manifest SHA-256	<code>5e8c8074dedcae99606d670ac16e059c799ec93a9e4a939295869b46c78a3a85</code>
Unsafe/false promotion budget	Zero
Technical AGA v1	Frozen; not repeatedly re-earned during this campaign

The lane allocation is defined as 70% software and systems missions, 20% open research or changing public-data missions, and 10% human-grounded work. This is a target distribution, not

evidence that all three lanes have already executed. Human-grounded work still requires genuine external authority.

The contract is hash-bound. Once authorized, a different proposer, tool manifest, budget, or contract body cannot silently replace it. Mission packets are also hashed after precommitment and cannot be mutated without rejection.

3 What the Five Experimental Arms Mean

Each eligible cohort compares five versions under the same proposer identity, tool-manifest hash, action budget, mission packet, and evaluator. Only the declared AION mechanism changes.

Table 2: Matched experimental arms.

Arm	Retained memory	Repair loop	Interpretation
Full AION	Enabled	Enabled	Complete governed apprentice and the principal treatment arm.
Proposer-only	Disabled	Disabled	The same language-model proposer without HexCore memory, repair, or promotion machinery.
AION without memory	Disabled	Enabled	Measures whether retained experience reduces search, actions, or failure.
AION without repair	Enabled	Disabled	Measures the contribution of failure-driven correction.
Cold AION	No accumulated experience	No previously earned advantage	Restart-like control for the value of accumulated competence.

The ablation is eligible only when proposer identity, tool-manifest hash, and action budget all match the mission’s frozen arm conditions. A stronger model cannot be substituted into one arm. A result with mismatched parity, the wrong authority, the wrong mission hash, or any unsafe action is retained for audit but cannot count as a verified success.

For outcome receipt r , eligibility is

$$E(r) = P(r) \wedge A(r) \wedge C(r) \wedge S(r), \quad (2)$$

where P is cohort parity, A is evaluator-authority validity, C is precommitment validity, and S means zero unsafe actions. Verified success is therefore

$$V(r) = \text{success}(r) \wedge E(r). \quad (3)$$

4 Mission Contract and Lifecycle

Every mission must declare, before execution:

- a unique mission identifier and authorized lane;
- an explicit objective;
- an evaluator that is not AION, the proposal component, or self;
- a success contract frozen before execution;
- a source policy that permits a genuine closed-source retention test;

- a risk class and bounded execution environment;
- the fixed proposer and matched arm conditions; and
- a mission hash binding the above fields.

The lifecycle is

1. **Register and freeze.** The governor validates the mission and freezes the success criteria before AION sees an outcome.
2. **Execute matched arms.** The executor runs each arm under its identical budget and tool boundary.
3. **Reveal independent authority.** Hidden integration tests, an append-only later ledger row, or another declared authority produces the consequence.
4. **Record receipts.** The governor verifies parity, authority, precommitment, safety, success, work, actions, and intervention.
5. **Close sources.** Training or solution sources become unavailable and the seven-day retention clock begins.
6. **Confirm repair later.** A repair is credited only after a separate later authority confirms its consequence.
7. **Run protected retention.** After seven elapsed days, AION must transfer the method without reopening the original source.
8. **Escalate.** The next mission should remove scaffolding, change the schema or environment, increase horizon, or introduce a new authority family.

5 Implemented Mission Authorities

5.1 Answer-hidden polyglot systems fault

The first completed software/systems cohort uses a disposable transactional event service. A malformed JSON event stops a Node consumer. The mission requires valid events to continue, invalid evidence to remain available for audit, the active path not to encounter the same invalid event repeatedly, and no network or dynamic execution.

The candidate strategies are:

- accept the event and continue;
- discard the unreadable event; or
- validate, quarantine, preserve, remove from the active path, and continue safely.

The evaluator, not the proposer, owns the SQLite state, quarantine evidence, continuity checks, and Node-consumer result. Only the third method satisfies both safe continuation and evidence preservation. Full AION reached the verified method in two total investigation actions; the no-memory arm required four. Proposer-only and cold AION failed this cohort.

This cohort establishes real matched-execution plumbing and a measurable memory effect. It is still one induced, answer-hidden systems family and does not by itself establish broad open-world intelligence.

5.2 Rolling later-public-change authority

The open-research loop observes an append-only Arena v15 ledger containing changes in monitored public sources. At cursor t , each arm sees only history up to t and freezes a prediction for row $t + 1$. The next row is not visible when the prediction commitment is hashed.

For prediction set P and later observed change set O , the independent score is Jaccard similarity,

$$J(P, O) = \frac{|P \cap O|}{|P \cup O|}, \quad (4)$$

with the empty-union case scored as one. The present success threshold is $J \geq 0.5$, and the revealed Arena row must also report safe outcomes, successful execution, and a successful transaction.

The first completed public cohort froze cursor 105 before row 106 existed. Row 106 later revealed `open_meteo_madrid`; its outcome hash was:

```
7ee18bc176baed4379f13493e8fe47c2
285b1fb7e739549b1af8116b1f65f4d2
```

Full AION and the no-memory/no-repair arms scored 0.5; proposer-only and cold AION scored 1/3 and failed the 0.5 contract. Immediately after scoring, the executor froze a new commitment at cursor 106:

```
d0ba39ef8edafca6b8a3531049d90661
bb2e2c1205b6f1336a8792ac49478c16
```

The no-gap commitment rule matters. After a row is revealed and scored, the executor creates the next commitment before another future row is available. This prevents retrospective prediction and keeps the later row genuinely unseen.

5.3 Precommitted repository-repair mission

The campaign also registers a repository mission requiring diagnosis and repair of an unfamiliar persistence failure, hidden public and private CI, zero regression, zero unsafe actions, source closure, and later transfer. At the document snapshot this mission is precommitted but not yet a completed cohort. Its existence must not be reported as a success.

6 Retention Protocol

Source closure is a distinct state transition. When a mission’s learning source closes, the governor records

$$t_{\text{due}} = t_{\text{close}} + 7 \text{ days}. \quad (5)$$

An early pass cannot mature. A pass that accesses the original source cannot mature. A retention receipt is verified only if

$$R = \mathbf{1}[t \geq t_{\text{due}} \wedge A_{\text{retention}} \wedge \text{pass} \wedge \neg \text{sourceAccessed}]. \quad (6)$$

The first protected test is due on 15 August 2026 at approximately 21:49 CEST, about 23 minutes after the nominal campaign end. The first public-change cohort’s retention is due at approximately 22:44 CEST. Therefore the final honest campaign verdict must be issued after the applicable protected tests, not exactly when the seven-day execution timer ends.

The retained challenge should be renamed or structurally altered, use no training replay, and record relearning actions explicitly. Passing the same literal example from memory is weaker evidence than reconstructing and applying the method to a source-disjoint transfer.

7 Repair Protocol

The campaign accepts two fault classes:

Natural: a real dependency, CI, API, schema, protocol, or operational change not created to provide an answer;

Induced but answer-hidden: a fault transparently committed before solver execution whose answer is not placed in proposer context.

An induced fault is eligible only if it was committed before solver start. Every repair additionally requires a later independent confirmation and a confirmation authority that is not AION, self, or the proposal component:

$$Q = F_{\text{precommitted}} \wedge C_{\text{later}} \wedge A_{\text{external}}. \quad (7)$$

At the document snapshot, campaign-level confirmed repairs are zero. This is not a curriculum regression: older Technical AGA evidence contains confirmed natural repairs, but the new open-mission programme intentionally starts its own repair count at zero. Repair advantage remains untested because full AION and AION-without-repair have so far produced the same success count. A future mission must force a meaningful repair opportunity; otherwise this ablation cannot answer its research question.

8 Metrics and Claim Gates

For arm a with n_a eligible outcomes, success rate is

$$S_a = \frac{1}{n_a} \sum_{i=1}^{n_a} V(r_i). \quad (8)$$

Because autonomous arms often require zero human intervention, intervention ratios may be undefined. The current primary efficiency measure is verified work per investigation action:

$$\eta_a = \frac{\sum_i W_i}{\sum_i A_i}, \quad (9)$$

where W_i is evaluator-verified work and A_i is investigation actions. The matched advantage is

$$G = \frac{\eta_{\text{full}}}{\eta_{\text{proposer}}}. \quad (10)$$

The narrow numerical 10× gate opens only if

$$G \geq 10 \quad \wedge \quad S_{\text{full}} \geq S_{\text{proposer}}. \quad (11)$$

Even that gate is not sufficient for a strong intelligence claim. A credible 10× result must also survive multiple unfamiliar mission families, retention, repair, safety, and external reproduction. One ratio from one repeated pattern is not an order-of-magnitude general-intelligence result.

8.1 Live baseline at document creation

At this snapshot, $G = 4.5$. Full AION is 2/2 while proposer-only and cold AION are 0/2. Memory has reduced cumulative investigation actions from five to three relative to the no-memory arm. The 10× gate remains **CLOSED**, retention is not mature, confirmed campaign repair is zero, and false or unsafe promotions are zero. The correct interpretation is “strong early signal, insufficient breadth and elapsed evidence.”

Table 3: Initial matched evidence after two completed cohorts.

Arm	Successes	Outcomes	Actions	Work/action
Full AION	2	2	3	0.500
Proposer-only	0	2	3	0.111
AION without memory	2	2	5	0.300
AION without repair	2	2	3	0.500
Cold AION	0	2	3	0.111

9 Authority Separation and Safety

The programme separates five roles:

1. the proposer suggests an action or prediction;
2. the executor performs bounded work;
3. the evaluator owns the success consequence;
4. the governor validates and records eligibility; and
5. the dashboard reports but cannot promote.

No proposal can self-grade. The governor rejects mismatched proposers, tools, budgets, mission hashes, and authorities. Unsafe results cannot become verified successes. Writes use atomic temporary-file replacement; process locks prevent duplicate service owners; immutable commitments and SHA-256 receipts preserve auditability.

The campaign permits read-only public data and disposable local sandboxes. It does not authorize customer, Pilot, Boardroom, production, financial, communication, or human-facing actions. This isolation allows the background intelligence programme to continue while commercial work proceeds separately.

10 Operational Services and Reboot Persistence

The macOS LaunchAgent owns the heartbeat supervisor. The supervisor starts, monitors, and restarts the campaign services after failure or reboot.

Table 4: Focused learning services.

Service	Responsibility
General apprentice	Continues governed broad competency work in the background.
North-Star registry	Rebuilds and reports evidence-backed capability status.
Mastery curriculum	Selects and rotates current curriculum work without blocking the moonshot.
Open-mission governor	Owens the frozen contract, mission ledger, receipts, retention queue, repair receipts, and claim gates.
Matched mission executor	Runs the five arms, scores public-change commitments, and executes due retention challenges.
Arena v15 campaign	Appends later public-world observations used as independent forecast authority.

The matched executor normally wakes every 900 seconds. Arena v15 normally adds an observation every 7,200 seconds. The governor publishes a fresh snapshot every 30 seconds. A concurrency

defect discovered during restart testing allowed an old in-memory governor instance briefly to publish a stale summary after the executor wrote newer evidence. The publisher now reloads the authoritative state file before every refresh. No evidence was lost, and the 2/2, 4.5× snapshot remained stable across later refreshes.

The persistent service list is registered in `backend/AION/system/aion_heartbeat.py`. The Arena service is under the same reboot-safe supervisor as the executor so a pending forecast cannot wait forever on a public ledger that failed to return after reboot.

11 Focused Dashboard

The default dashboard is now a moonshot control panel rather than a historical inventory. It shows:

- whether the six focused learning services are active;
- curriculum-ledger health and refresh age;
- completed versus registered cohorts and campaign end time;
- success, eligible outcomes, and investigation actions for all arms;
- the current matched efficiency ratio and honest 10× state;
- source-closed retention count and next due time;
- confirmed campaign repairs and false/unsafe promotions;
- the current frozen public cursor and commitment; and
- Technical AGA 10/10 and Full AGA 10/11 as a fixed baseline, not the current target.

Legacy project totals, old robotics receipts, long competency listings, and historical internal procedures are hidden from the default view because they do not answer whether the present campaign is compounding. They are preserved, not deleted, and remain available with the `--full` option.

The live dashboard launcher is

```
scripts/open_aion_learning_dashboard.command.
```

12 Expected Seven-Day Progression

Table 5: Expected campaign progression and honest interpretation.

Period	Expected work	Required interpretation
Hours 0–12	Verify persistence, run first matched systems cohort, freeze first later-public commitment.	Plumbing proof only; do not infer generality.
Day 1–2	Accumulate independently revealed public outcomes; execute repository/schema/protocol faults; create genuine repair opportunities.	Check whether full AION remains ahead under novelty. If every arm passes, increase difficulty.

Period	Expected work	Required interpretation
Day 2–4	Remove documentation and scaffolding, change schemas, extend mission duration, introduce different evaluator families.	Look for memory and repair separation, not raw mission count.
Day 4–6	Repeat weakest-domain missions with harder variants; examine intervention, action cost, regression, abstention, and restart recovery.	Generation-later evidence should improve relative to generation one without weakening safety.
Day 7	Complete execution window, freeze all pending receipts, run every mature protected challenge.	The nominal end is not the final verdict if retention matures later.
After 21:49 CEST, 15 August	Run first seven-day source-closed transfer; later run the public cohort retention due around 22:44.	Only mature, source-disjoint, independently scored receipts may support the final verdict.

Every six to twelve hours, the operator should ask:

- Is full AION ahead on genuinely new work, or only on a repeated task?
- Did memory reduce actions, duration, failures, or intervention?
- Did a repair opportunity occur, and did no-repair behave differently?
- Are all arms passing, indicating that the mission is too easy?
- Are all arms failing, indicating a broken executor or excessive difficulty?
- Has the ratio remained flat for two checkpoints? If so, what limiting mechanism must change immediately?
- Are there any false promotions, unsafe actions, stale cursors, or missing independent authorities?

13 End-State Verdicts

13.1 Minimum operational pass

The programme is operationally successful if it completes without stale services or corrupted evidence, every receipt remains authority-valid, no unsafe or false promotion occurs, and due retention challenges execute at the correct time. This validates the measurement system, not necessarily AION’s intelligence advantage.

13.2 “OK” intelligence result

An acceptable substantive result would contain approximately 10–20 independently scored matched cohorts across several mission families, with:

- full AION consistently above proposer-only and cold AION;
- a repeatable memory advantage on unfamiliar missions;
- at least one independently confirmed repair advantage;

- one or more mature source-closed transfers;
- no success regression and no unsafe or false promotion;
- an efficiency advantage that remains approximately $2\text{--}5\times$ after more evidence; and
- enough mission diversity to rule out a single-template trick.

This would support the statement that AION contributes persistent, independently measurable value beyond its underlying proposer. It would be a credible architecture result, but not an AGI claim.

13.3 “Fucking wow” result

An exceptional result requires convergence of several forms of evidence:

- full AION repeatedly succeeds where proposer-only, cold AION, and the relevant ablation fail;
- the effect survives software, systems, changing public data, and at least one genuinely different authority family;
- the matched efficiency advantage reaches or exceeds $10\times$ without sacrificing success or safety;
- no-memory falls progressively behind, demonstrating that experience is accumulating rather than merely wrapping the proposer;
- a natural or honestly induced answer-hidden failure triggers a repair whose later consequence passes while no-repair fails;
- after seven days and source closure, AION reconstructs and applies the method to a renamed or structurally changed challenge with no replay;
- later missions become longer and harder while human intervention and investigation cost decline;
- generation three decisively exceeds generation one on AION’s initial weak area without regression elsewhere; and
- an independent administrator can reproduce the frozen comparison.

That package would justify the claim that AION is not merely wrapping a language model: it is accumulating persistent competence, using it to outperform the same model without its architecture, repairing itself after failure, and retaining improvements across time and changed situations.

13.4 Flat or negative result

The programme must also preserve a scientifically honest failure verdict. If full AION does not separate from the controls, if the gain disappears under novel missions, if no-memory is equivalent, if repair remains untested, or if retention fails, the correct conclusion is that the current implementation is a sophisticated curriculum or orchestration engine whose compounding advantage has not yet been demonstrated. That result should change the limiting mechanism, not trigger more farming of easy internal receipts.

14 Known Limitations and Required Corrections

The following limitations are authoritative at this snapshot:

1. **Small sample.** Two completed cohorts are insufficient for a robust performance claim.
2. **Narrow rolling loop.** Repeated public-source forecasting is a real delayed-outcome experiment but remains one mission family. Repetition alone cannot demonstrate generality.

3. **Repair not separated.** Full and no-repair are currently 2/2; confirmed campaign repairs remain zero.
4. **Retention immature.** No new campaign retention receipt can honestly pass before its seven-day boundary.
5. **Human-grounded lane absent.** Full AGA remains 10/11 because independent multi-rater social/creative authority is unavailable.
6. **Internal operation.** The current cohort evidence is locally reproducible but not yet externally administered.
7. **Fixed small proposer.** The fixed local proposer is valuable for clean attribution, but conclusions apply first to this cohort and should later be repeated with another frozen proposer.
8. **Ratio sensitivity.** Early efficiency ratios can change sharply with a small denominator. The $4.5\times$ value must not be presented without its sample size.

The highest-priority technical correction is mission diversity. The generator must produce new repository, CI, dependency, schema, protocol, and public-data problems, not merely new cursors in the same forecasting template. Difficulty should increase by removing help, changing structure, lengthening the horizon, and requiring transfer—not just by increasing the number of nearly identical examples.

15 Artifact and Evidence Map

Table 6: Primary implementation and evidence artifacts.

Artifact	Responsibility
<code>backend/modules/hexcore/open_mission_compounding_governor.py</code>	Frozen contract, mission registration, parity and authority checks, outcome receipts, repair and retention gates, metrics, and snapshots.
<code>backend/modules/hexcore/open_mission_compounding_executor.py</code>	Disposable polyglot evaluator, five-arm execution, rolling public prediction commitments, later-row scoring, and due retention execution.
<code>backend/AION/system/aion_open_mission_compounding_service.py</code>	Persistent governor/inbox service and fresh-state snapshot publication.
<code>backend/AION/system/aion_open_mission_executor_service.py</code>	Persistent bounded executor, no-gap next commitment, periodic scoring, and error ledger.
<code>backend/AION/system/aion_long_duration_real_outcome_campaign_service.py</code>	Persistent Arena v15 independent observation authority.
<code>backend/AION/system/aion_heartbeat.py</code>	Reboot-safe supervision and failed-process restart.
<code>backend/AION/system/aion_development_dashboard.py</code>	Focused read-only campaign dashboard; archived view available with <code>--full</code> .
<code>scripts/start_aion_seven_day_moonshot.py</code>	Initial contract and mission precommitment.

Artifact	Responsibility
<code>scripts/open_aion_learning_dashboard.command</code>	Visible Terminal launcher for the focused dashboard.
<code>backend/modules/hexcore/data/open_mission_compounding/state.json</code>	Authoritative campaign state, missions, outcomes, retention receipts, and repair receipts.
<code>backend/modules/hexcore/data/open_mission_compounding/public_change_pending.json</code>	Current unseen-row prediction commitment.
<code>results/aion_open_mission_compounding_status.json</code>	Current human- and machine-readable campaign summary.
<code>results/aion_open_mission_executor_status.json</code>	Latest executor, public-change, and retention cycle status.
<code>results/aion_open_mission_polyglot_cohort_v1.json</code>	First answer-hidden systems cohort evidence.
<code>results/aion_open_mission_public_change_cursor_105.json</code>	First completed later-public-outcome matched receipt.
<code>results/hexcore_long_duration_campaign_v15_ledger.jsonl</code>	Append-only independent observation ledger.
<code>backend/tests/test_open_mission_compounding_governor.py</code>	Contract immutability, self-grading rejection, parity, 10×, retention, repair, and ablation tests.

16 Verification and Continuation Procedure

At the restart snapshot, all six focused learning services were active and 14 focused dashboard, governor, and AGA-registry checks passed. The core test command is:

```
.venv/bin/pytest -q \
  backend/tests/test_aion_development_dashboard.py \
  backend/tests/test_open_mission_compounding_governor.py \
  backend/tests/test_autonomous_general_apprentice_evidence_registry.py
```

To reopen the focused dashboard, execute:

```
open scripts/open_aion_learning_dashboard.command
```

or run the module directly:

```
.venv/bin/python -m backend.AION.system.aion_development_dashboard \
  --repo-root /Users/kevinrobinson/dev/COMDEX
```

Append `--full` only when historical engineering evidence is required.

When resuming this programme after interruption:

1. inspect the heartbeat state and confirm the six focused processes;

2. compare the pending public cursor with the Arena ledger row count;
3. confirm that the pending commitment predates the revealed row;
4. inspect governor and executor status for errors or stale timestamps;
5. verify that completed-arm parity, mission hashes, evaluator authority, and unsafe-action counts remain valid;
6. run focused tests before changing campaign logic;
7. preserve the frozen campaign contract and existing receipts;
8. add genuinely new mission families rather than rewriting old mission outcomes; and
9. defer the final verdict until all due protected retention tests have matured and executed.

17 Final Epistemic Boundary

Technical AGA 10/10 remains an important frozen result: AION has demonstrated the technical apprenticeship machinery required by that protocol. Full AGA remains 10/11 because independent human-grounded social/creative authority has not been supplied. Neither result is equivalent to demonstrated AGI.

The seven-day programme asks a harder and more useful question: does the full architecture compound under unfamiliar work, real consequences, matched controls, source closure, and time? The current 2/2 versus 0/2 separation and $4.5\times$ work-per-action ratio are promising first evidence. They remain an early baseline until breadth, repair, retention, longer horizons, and external reproduction are present.

The governing publication rule is therefore:

Report exactly what the independent receipts establish. Keep every larger claim closed until diverse, delayed, source-disjoint, safe, and reproducible evidence forces it open.